



Introduction to Neuroevolution

CISC 455/851

Ryan Zhou

Outline

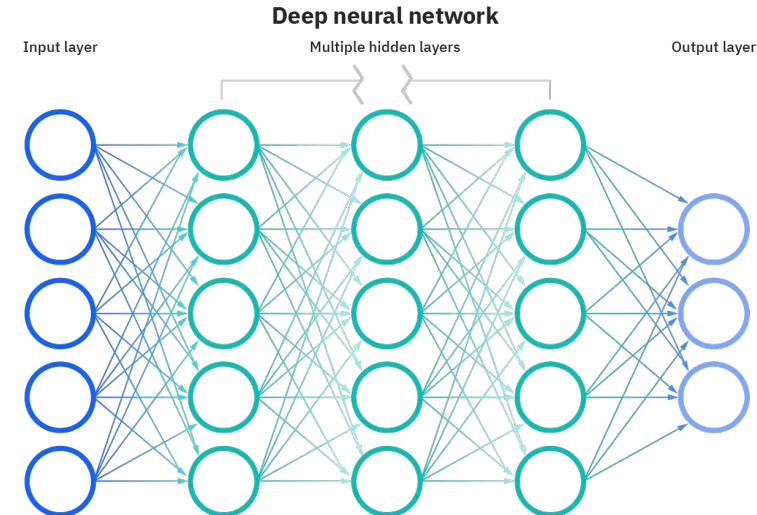


- What is neuroevolution?
- Classical neuroevolution
- Hybrid neuroevolution
- The future of neuroevolution

Neural networks



- Another type of machine learning model
- Network of “neurons” with connections between them
- Each connection has a **weight**
- Usually trained with gradient descent



What's special about neural networks?

- Neural networks can be huge ->
- Architecture matters, a lot
- Many different ways to train
- Biological inspiration

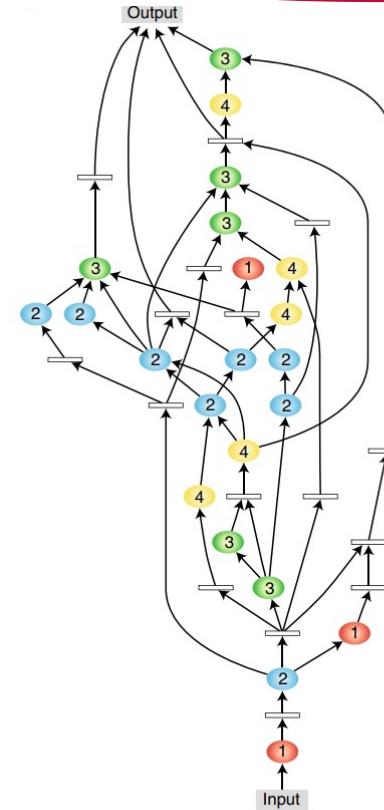


Name	Release date ^[a]	Developer	Number of parameters ^[b]	Corpus size
GPT-4	March 2023	OpenAI	Unknown ^[f]	Unknown
PaLM (Pathways Language Model)	April 2022	Google	540 billion ^[40]	768 billion tokens ^[39]
Minerva	June 2022	Google	540 billion ^[44]	38.5B tokens from webpages filtered for mathematical content and from papers submitted to the arXiv preprint server ^[44]
Megatron-Turing NLG	October 2021 ^[27]	Microsoft and Nvidia	530 billion ^[28]	338.6 billion tokens ^[28]
BERT	2018	Google	340 million ^[16]	3.3 billion words ^[16]
Gopher	December 2021	DeepMind	280 billion ^[35]	300 billion tokens ^[36]
Ernie 3.0 Titan	December 2021	Baidu	260 billion ^{[29][30]}	4 Tb
GPT-3	2020	OpenAI	175 billion ^[9]	499 billion tokens ^[20]

Image source: Wikipedia

What is neuroevolution?

- Use an evolutionary algorithm with a neural network
- Can replace gradient descent, or work in conjunction with it



Why use neuroevolution?



- Gradient descent requires **differentiability**
- Many aspects of a neural network are not differentiable
 - Architecture
 - Hyperparameters
 - Learning rules
 - Choice of objective
- Maintains a solution population
 - Enables extreme exploration and massive parallelization
- Good if there's lots of local minima that gradient descent could get stuck in



Classic Neuroevolution

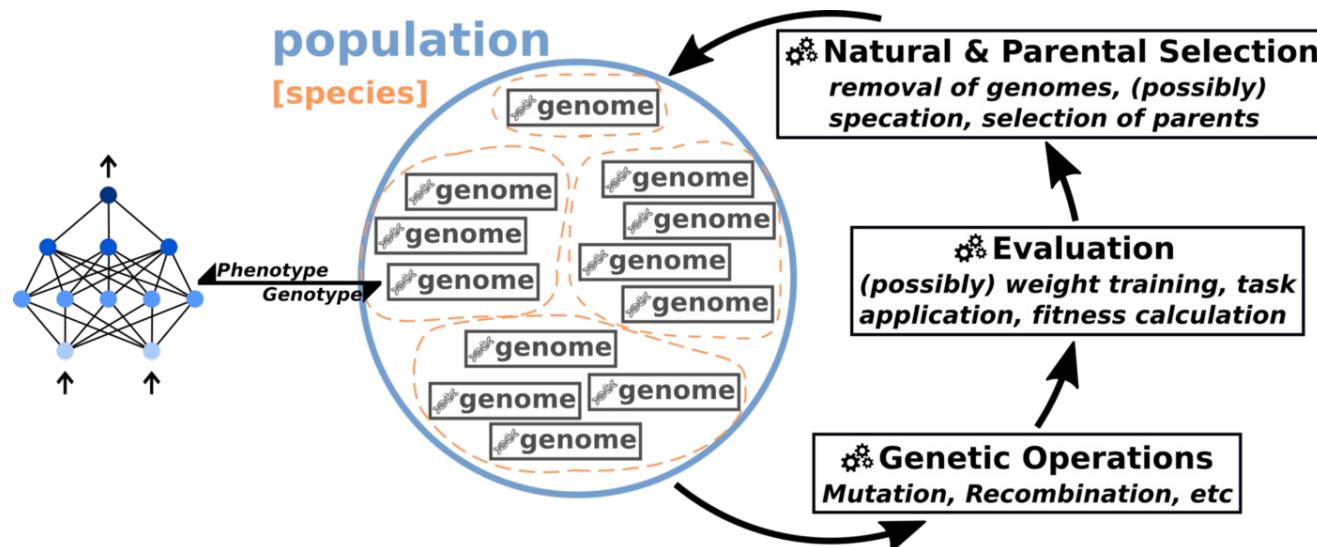
Training with evolution!

Classic Neuroevolution



Queen's
UNIVERSITY

- Use an evolutionary algorithm to evolve a network
- EA does all the work here

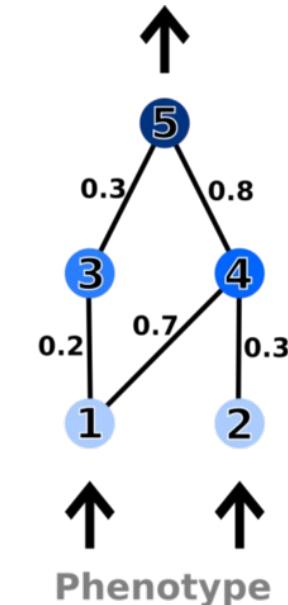


Encoding a Network

- Direct encoding
- 1 gene = 1 weight
- Architecture can be fixed/mutable
 - Easier to fix it and evolve weights only

genome			
gene: #1	in: 1	out: 3	weight: 0.2
gene: #2	in: 1	out: 4	weight: 0.7
gene: #3	in: 2	out: 4	weight: 0.3
gene: #4	in: 3	out: 5	weight: 0.3
gene: #5	in: 4	out: 5	weight: 0.8

Genotype



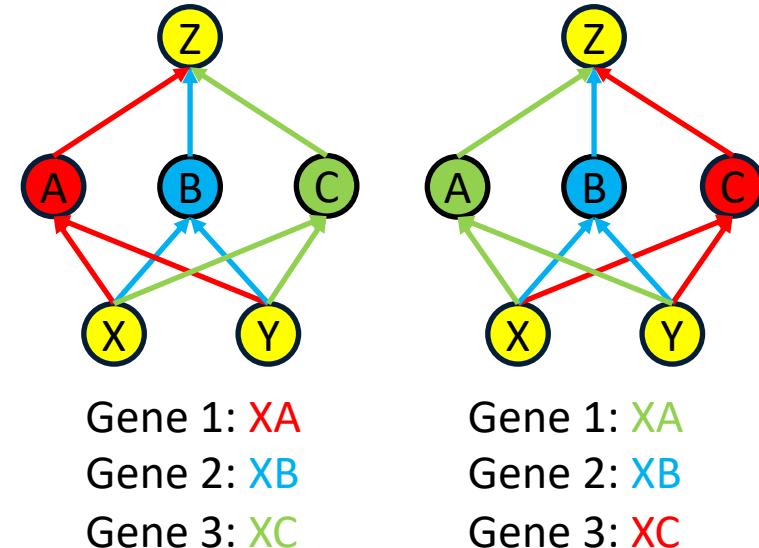
Basic Mutation



- Modify a weight
- That's it!
- Can use any of the real-valued mutation strategies

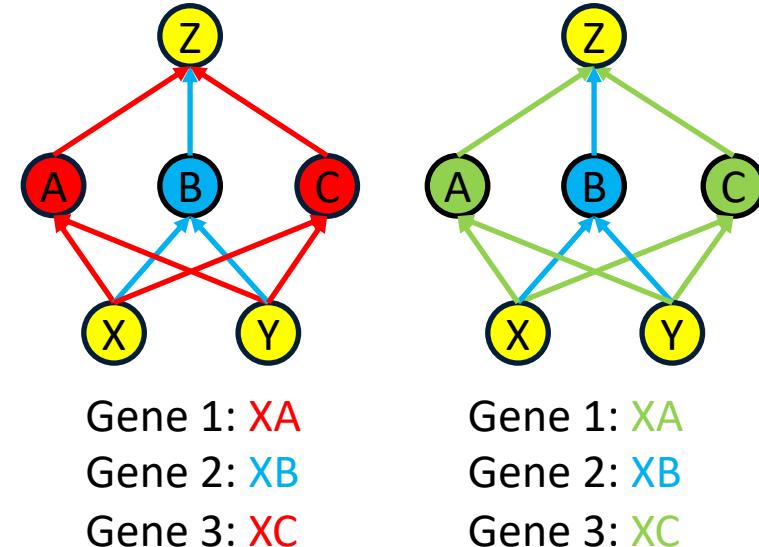
How do we do recombination?

- Main problem: networks can be permuted
- If you permute the order of the nodes, you'll get a functionally identical network...
 - But the genes will be in a different order



How do we do recombination?

- The parents were the same network
- But when you recombine them, you lose information
- Problem for many direct encodings
 - Each implementation will do recombination differently (or not at all)



How well does this work?



Queen's
UNIVERSITY

- Works well for problems where the correct solution isn't known
 - E.g. reinforcement learning type problems
- Reward signal can be sparse or misleading, so it makes sense to throw a lot of things at the wall and keep what sticks
- Direct encoding (1 gene = 1 weight) has some major drawbacks
 - Big networks need big genomes
 - Recombination is difficult

Applications

- Robotics
 - First running gait for the Aibo dog robot
 - Morphology of 3d printed robots (Lipson & Pollack 2000)
- Measurement of the top quark
- Agents for video games

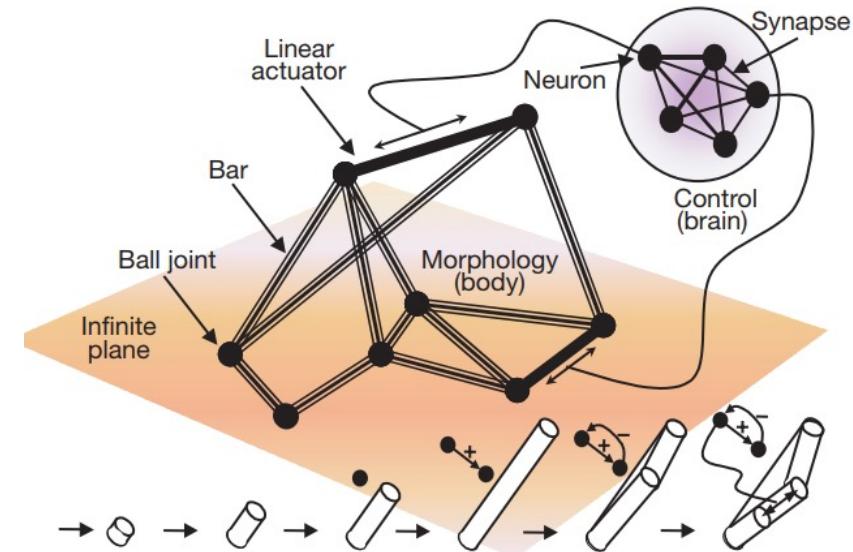
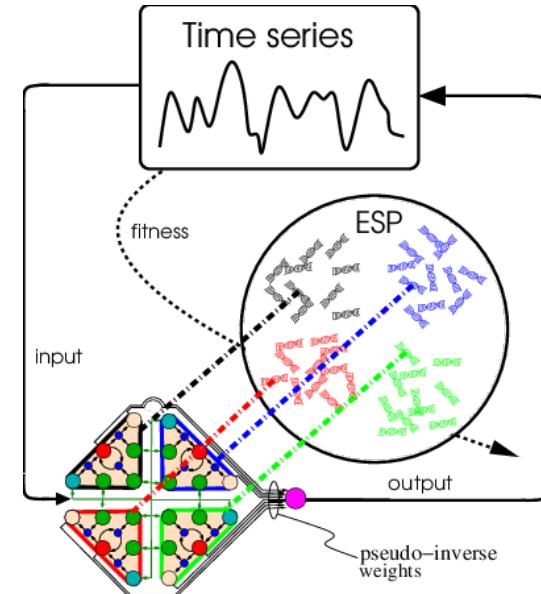


Image source: Lipson, H. and Pollack, J. Automatic design and manufacture of robotic lifeforms

Variation 1: Neurons are individuals



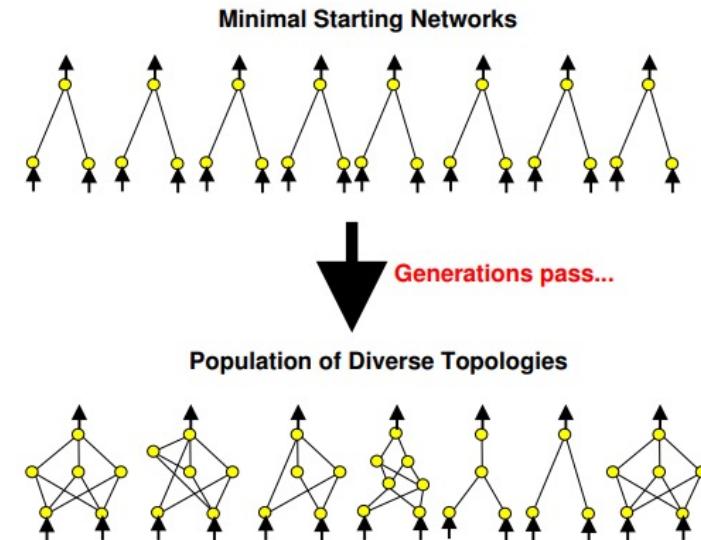
- Evolve individual neurons or small modules, and force them to work together
- Fitness is based on how well the whole network performs
- Recombination between neurons is easier than between full networks



Variation 2: Complexification



- It's hard to optimize a giant network right off the bat
 - Search space is too big
- Start small and build up
- Example: NEAT (Neuroevolution of Augmenting Topologies)



Variation 3: Indirect Encoding



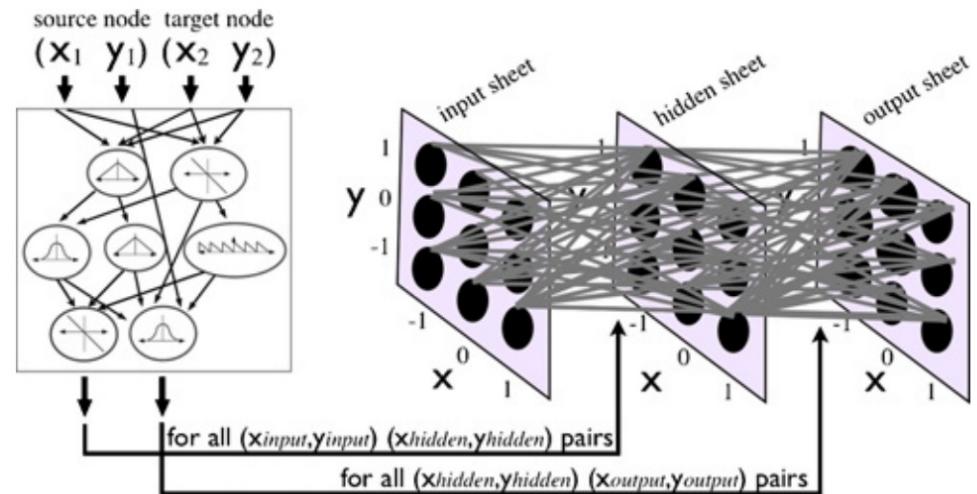
- 1 gene = 1 weight is inefficient
- Why not have 1 gene = multiple weights?
- Not all the connections in the brain are stored in DNA
 - Information is highly compressed
- Encode a set of rules to generate the weights, and evolve the rules

HyperNEAT



Queen's
UNIVERSITY

- Use a small network to describe the **pattern** of weights in the full network
- Kind of similar to attention
 - which can also be thought of as indirect encoding
- Adaptive HyperNEAT
 - Also generates an update rule so it can learn





Hybrid Neuroevolution

Designing with evolution

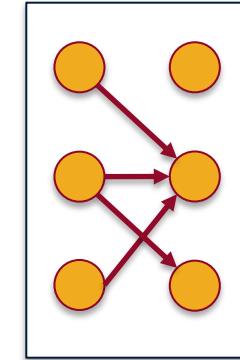
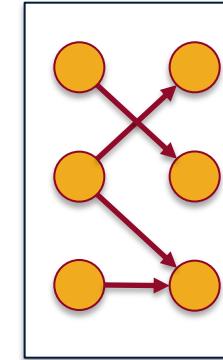
What is hybrid neuroevolution?



- Combine deep learning with evolution
- Use an EA to design the network, and train what you can
 - Initialization
 - Hyperparameters
 - Architecture
- Learning can actually speed up evolution

The Baldwin Effect

- Imagine an organism with a neural net brain
- There is only one correct way to connect the net so that it actually works
- How can evolution discover this?
 - Unless it hits on the exact right solution, there's no selection pressure
- “How Learning Can Guide Evolution”
 - Hinton and Nowlan



10/10

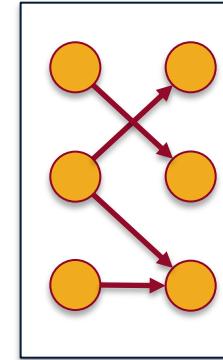
0/10

Pure evolution

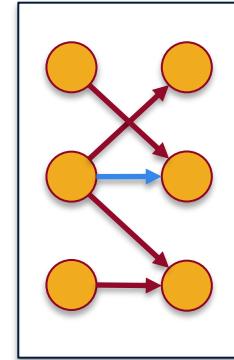


Queen's
UNIVERSITY

- With no learning, the neural net is fixed by the genome and never modified after
- Exact solution must be encoded in the genome
- No partial credit!
- It takes a long time to evolve the right solution since there's little selection pressure



10/10



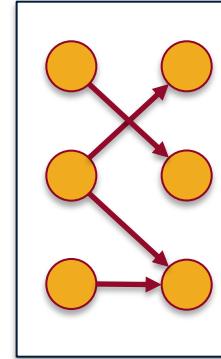
0/10

What if we allow them to learn?

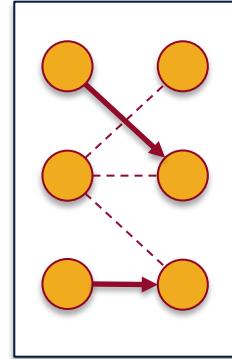


Queen's
UNIVERSITY

- Some of the connections are still fixed by the genome
- Others can be modified (“learnable”)
 - An organism can try out a certain number of configurations during its lifetime
 - Even if it doesn’t start with the right configuration, it can get partial credit if it finds the right one



10/10



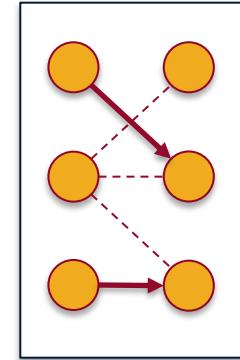
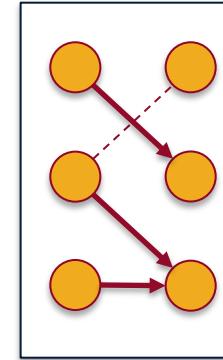
?/10

Observations



Queen's
UNIVERSITY

- If the genome gets some of the connections right, the organism is more likely to hit the correct solution during its lifetime
- The more correct connections we have prespecified, the faster this can happen
- Now we have selection pressure!
- Lesson: learning lets evolution explore faster



9/10

5/10

Neural architecture search

- Deep learning architectures are huge
- Use evolution to search over hyperparameters and architecture (meta EA)
- Genome defines the layers and sizes, and their hyperparameters

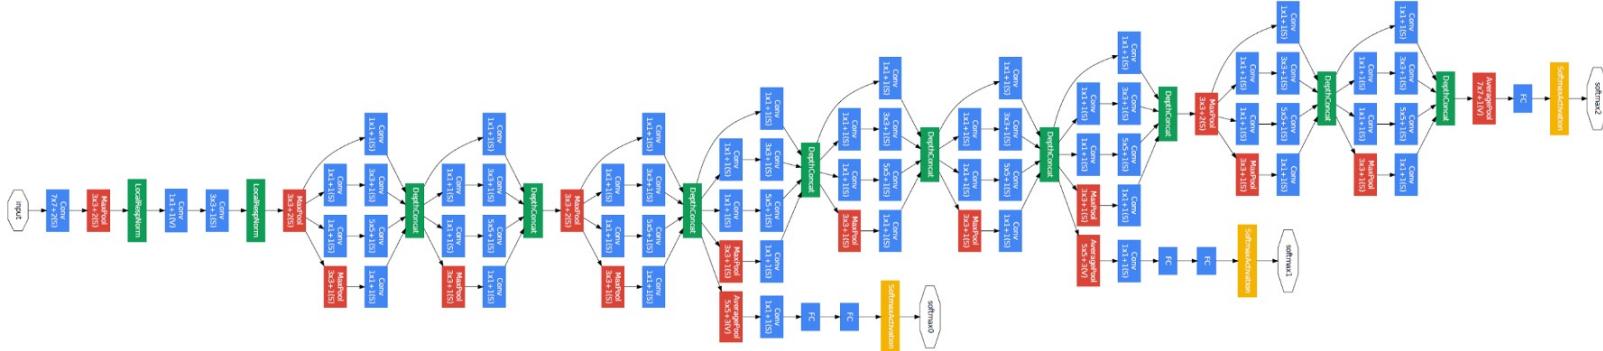


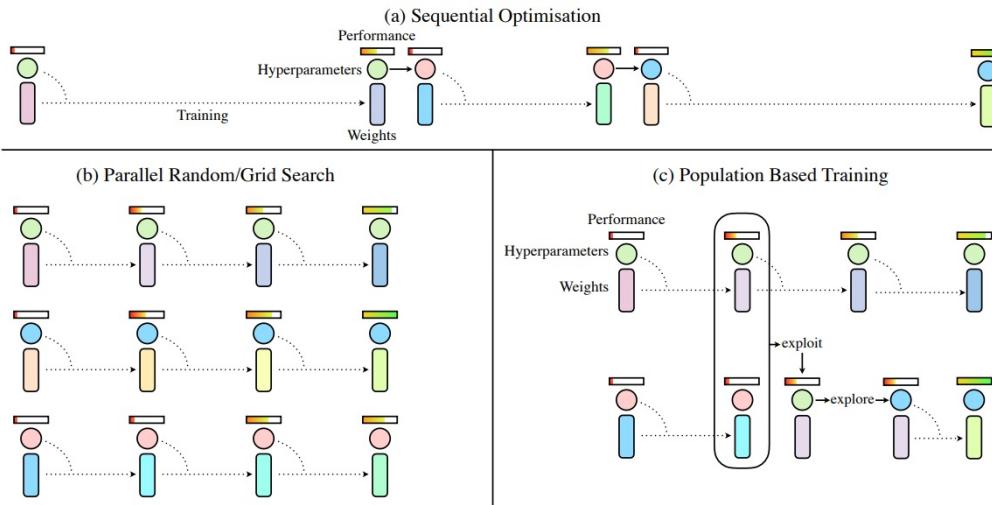
Image source: Szegedy et al. Going Deeper with Convolutions.

Hyperparameters: Population-based Training

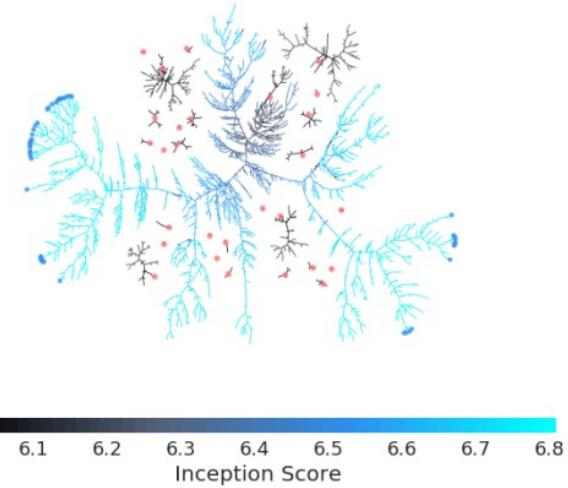


Queen's
UNIVERSITY

- Replace bad hyperparameters with mutations of good ones



GAN population development



Architecture: CoDeepNEAT



- Evolve small modules, and then evolve how to link up those modules

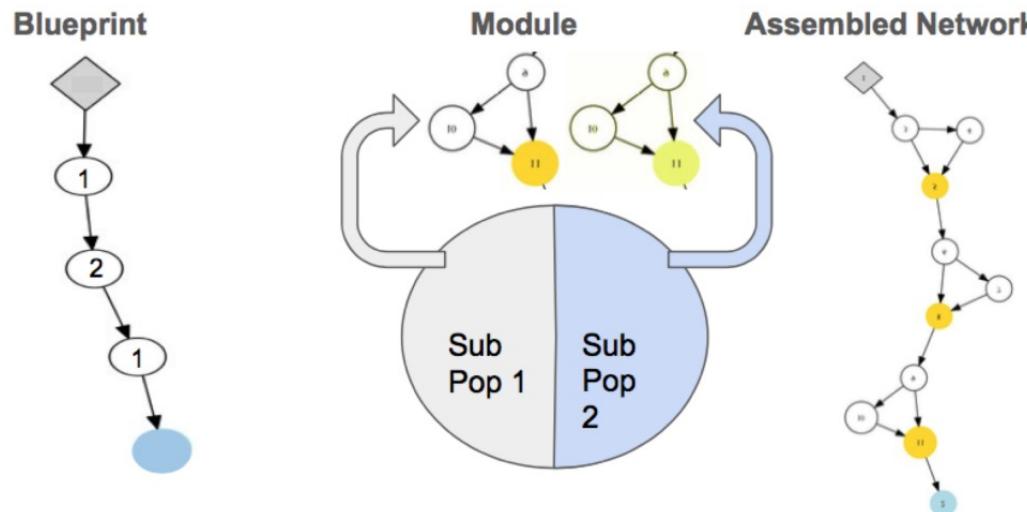
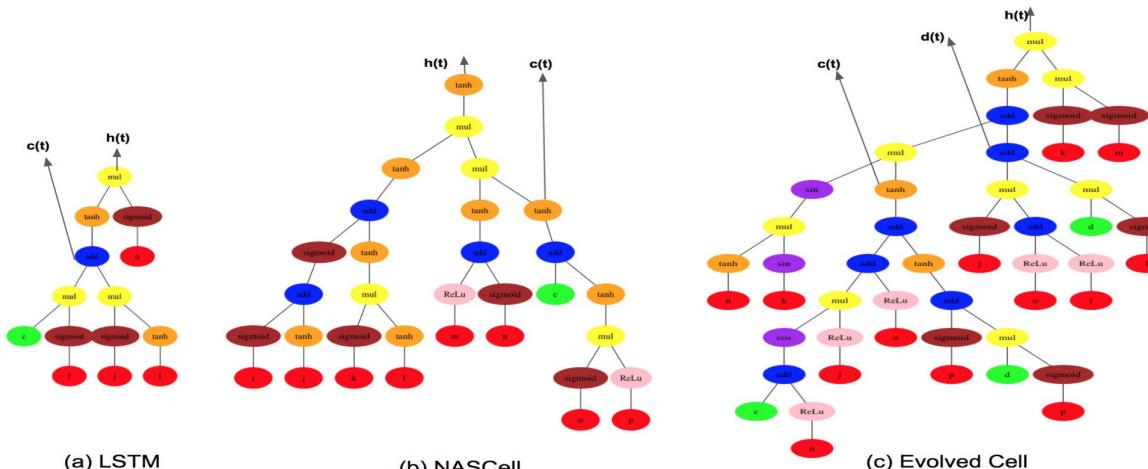


Image source: Miikkulainen et al. Evolving Deep Neural Networks.

Evolving neural network layers

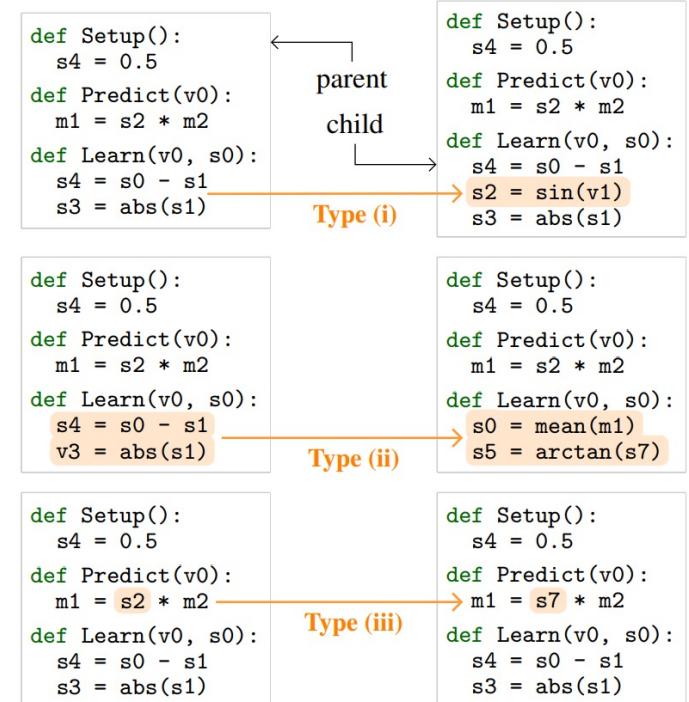
- Neural network layers are just programs
 - So why not apply genetic programming to evolve a better layer?
- Rawal and Miikkulainen: Evolve alternatives to LSTM



AutoML-Zero



- Evolve an entire machine learning algorithm from scratch
- Evolve three functions: setup, predict, learn
- Was able to discover backprop, dropout, weight averaging etc.
- Also done for RL (Co-Reyes et al)



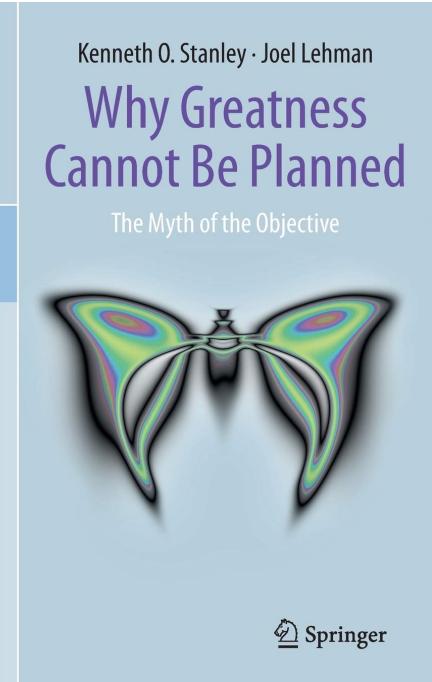


The Future of Neuroevolution

Open-ended Learning



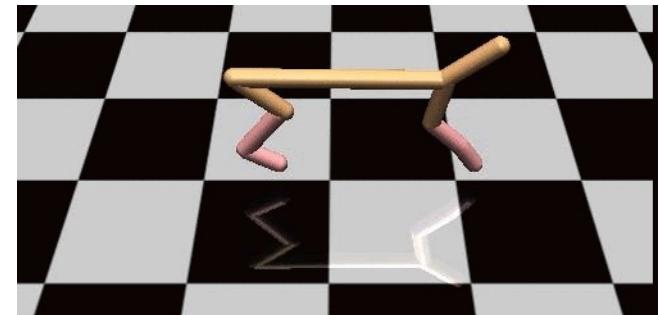
- Up until now, we used learning to come up with solutions for specific problems
- Can an algorithm generate problems for itself to solve?
- Open-ended behaviour
 - If we ran an algorithm for a billion years, would it still produce interesting results?



Diversity is all you need

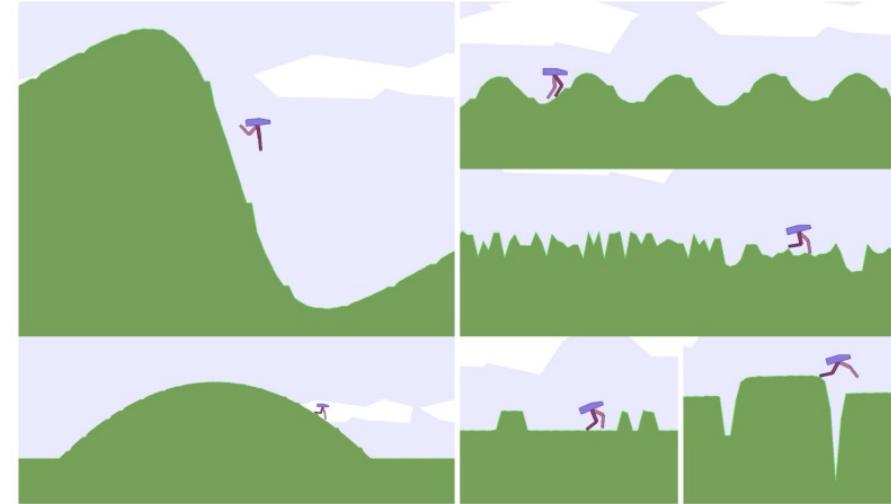


- Novelty search: only look for novel behaviours
- We don't care at all about performance
- Can learn to walk without being instructed to do so





- Paired open ended trailblazer (POET)
- Evolves a population of walking agents, and also increasingly harder environments
- Tries to make environments that are “interesting”
 - If an environment changes the ranking of the different walkers, it’s probably interesting



Evolution with large models



Queen's
UNIVERSITY

- Deep learning models are extremely capable now
- Can manipulate anything that can be represented as text
 - Code
 - Math
 - Molecules
 - DNA
 - Songs
- Rather than using evolution to guide deep learning, can deep learning guide evolution?

Evolutionary operators with LLMs



Queen's
UNIVERSITY

- Basic evolutionary operators
 - Mutation
 - Recombination
 - Selection
- How do we do this with a slogan? A poem? A program for drawing pictures?

Evolutionary operators with LLMs



Queen's
UNIVERSITY

- Just ask ChatGPT to do it!

The image shows a screenshot of a ChatGPT-like interface. On the left, there is a green square icon with a white AI symbol. To its right is a numbered list of five ice cream slogans. Below this list is a user instruction. At the bottom, another green square icon with an AI symbol is followed by a sixth slogan.

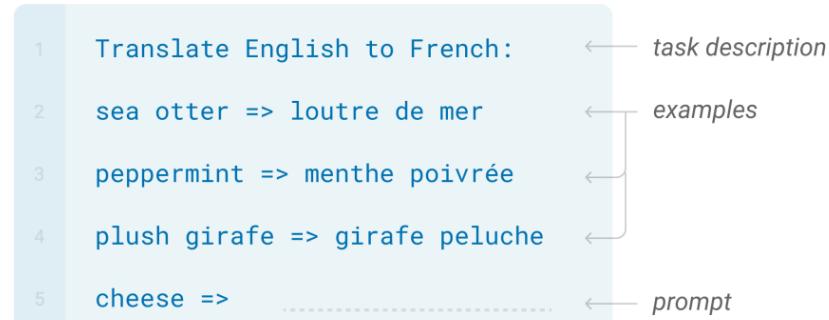
1. "Scoops of happiness in every cone!"
2. "Cool off with our delicious ice cream flavors!"
3. "Indulge in our creamy, dreamy treats!"
4. "Life is uncertain, but ice cream is a sure thing!"
5. "Taste the difference with our premium ice cream!"

RY Combine slogans 1 and 4.

6. "Certainty of happiness in every scoop!"

Why do it this way?

- “Smart” mutation and crossover
- Produces grammatically correct offspring
- Can understand complex requirements and goals
- Operator can be trained or even adapt on the fly with few-shot learning



“Smart” genetic programming



- Example: writing a function that draws a yellow circle
- Code LLM can't do it normally, but it can with evolution!
- How many generations would you need without smart crossover/mutation?



```
import math
import numpy as np

# Fixed version of draw()
def draw():
    """
    Draws a yellow circle with radius 10 in the middle of a 32 by 32 black image.

    Returns:
        np.ndarray: the image
    """
    pic = np.zeros((32, 32, 3))
    pic[16][16] = 1

    # Circle in middle of image
    for x in range(0, 32):
        for y in range(0, 32):
            distance = math.sqrt(math.pow(x - 16, 2) + math.pow(y - 16, 2))
            if distance > 10:
                continue

            pic[x][y] = [255, 255, 0]
    return pic
```



Neuroevolution is still evolving!

Questions or want to work on this? → ryan.zhou@queensu.ca